


Lecture 4_Memory

Pixel Art

1. Pixels are squares, individual dots, of color that are arranged on an up-down, left-right grid.
2. You can imagine an image as a map of bits, where zeros represent black and ones represent white.

```
1 1 0 0 0 0 1 1
1 0 1 1 1 1 0 1
0 1 0 1 1 0 1 0
0 1 1 1 1 1 1 0
0 1 0 1 1 0 1 0
0 1 1 0 0 1 1 0
1 0 1 1 1 1 0 1
1 1 0 0 0 0 1 1
```



3. *RGB*, or *red*, *green*, *blue*, are numbers that represent the amount of each of these colors.

Memory

1. We can interpret the whole Memory as concurrent blocks. And each smallest block represents a byte (8 bits).
2. The C language has two powerful operators that relate to memory:

`&` Provides the address of something stored in memory.

`&n` can be literally translated as "the address of n"

- Instructs the compiler to go to a location in memory.
*n as "the value stored in the address that n stores"

Example:

```
// Prints an integer's address
#include <stdio.h>

int main(void)
{
    int n = 50;
    printf("%p\n", &n);
}
```

The `%p` allows us to view the address of a location in memory. Here, it tells the `printf` to go to the address of `n`, and then print out the address.

Pointer

1. A *pointer* is a variable that stores the address of something. In other words, a pointer is an address in the computer's memory.
2. A pointer is usually stored as an 8-byte value

Use of &

```
int n = 50;
int *p = &n;
```

Here, `p` is a pointer that points to an `int`, so `p` stores the address of the `int` (`p` is the address of the `int`), and the `int` is `n`. The right side means the address of `n`. So the expression means, `p` is the address(points to the address) of `int n`.

Now leverage `p`, we can change the code into as followed:

```
// Stores and prints an integer's address
#include <stdio.h>
int main(void)
{
    int n = 50;
    int *p = &n;
    printf("%p\n", p);
}
```

Here, `%p` requires a data type of address, and `p` as a pointer satisfies. So no

more & is needed in indicating the variable is an address.

Use of *

```
// Stores and prints an integer via its address

#include <stdio.h>

int main(void)
{
    int n = 50;
    int *p = &n;
    printf("%i\n", *p);
}
```

Here, p pertains to be the pointer pointing to an int, and *p means go to the address and find the value, i.e. go to the address of n and find the value stored in that address, that is n. so *p is the value stored in the address that p stores.

Strings

1. A string is an array of characters ended with `\0`
2. In C, there is no data type called string, instead, there is *char**, which means a string is literally a pointer that points to a character. This make sense because a pointer tells the compiler where the first byte of the string exists in memory.
3. A string of characters is simply an array of characters identified by the location of its first byte.

```
// Declares a string with CS50 Library

#include <cs50.h>
#include <stdio.h>

int main(void)
{
    string s = "HI!";
    printf("%s\n", s);
}
```

is equal to

```
// Declares a string without CS50 Library

#include <stdio.h>

int main(void)
{
    char *s = "HI!";
    printf("%s\n", s);
}
```

This is what happens under the hood:

The cs50 library includes a struct as follows: `typedef char *string`

4. One cannot compare two strings using the `==` operator, since it is comparing the starting address of two strings, which are apparently different.

Copying and Malloc

String Copying

```
// Capitalizes a string

#include <cs50.h>
#include <ctype.h>
#include <stdio.h>
#include <string.h>

int main(void)
{
    // Get a string
    string s = get_string("s: ");


    // Copy string's address
    string t = s;

    // Capitalize first letter in string
    t[0] = toupper(t[0]);

    // Print string twice
    printf("s: %s\n", s);
    printf("t: %s\n", t);
}
```

Even if we only change the first character of `t`, the first letter of `s` and `t` will both be capitalized, given that `s` and `t` points to the same string address because no memory is allocated to `t`

You can visualize the above code as follows:



Malloc and Free

1. `Malloc` allows you to allocate a block of a specific size of memory.
2. `Free` allows you to tell the compiler to *free up* that block of memory you previously allocated.
3. One way to copy an authentic copy of a string is as followed:

```
// Capitalizes a copy of a string

#include <cs50.h>
#include <ctype.h>
#include <stdio.h>
#include <stdlib.h>
#include <string.h>

int main(void)
{
    // Get a string
    char *s = get_string("s: ");
```

```

// Allocate memory for another string
char *t = malloc(strlen(s) + 1);

// Copy string into memory, including '\0'
for (int i = 0; i <= strlen(s); i++)
{
    t[i] = s[i];
}

// Capitalize copy
t[0] = toupper(t[0]);

// Print strings
printf("s: %s\n", s);
printf("t: %s\n", t);
}

```

Notice that `malloc(strlen(s) + 1)` creates a block of memory that is the length of the string `s` plus one. This allows for the inclusion of the null `\0` character in our final copied string.

One way to optimize the code is to define `n = strlen(s)` in the left-hand side of the for loop, so that it won't calculate `strlen(s)` in every loop.

```

for (int i = 0, n = strlen(s); i <= n; i++)
{
    t[i] = s[i];
}

```

Good news is that C has a built-in function to copy strings called `strcpy`, so we can directly malloc memory, then call the `strcpy` function.

```

// Capitalizes a copy of a string using strcpy

#include <cs50.h>
#include <ctype.h>
#include <stdio.h>
#include <stdlib.h>
#include <string.h>

int main(void)
{

```

```

// Get a string
char *s = get_string("s: ");

// Allocate memory for another string
char *t = malloc(strlen(s) + 1);

// Copy string into memory
strcpy(t, s);

// Capitalize copy
t[0] = toupper(t[0]);

// Print strings
printf("s: %s\n", s);
printf("t: %s\n", t);
}

```

4. If something goes wrong in `Malloc` , it returns `NULL` , which can be used to check the validation of Malloc operation

```

// Allocate memory for another string
char *t = malloc(strlen(s) + 1);
if (t == NULL)
{
    return 1;
}

```

Valgrind

1. *Valgrind* is a tool that can check to see if there are memory-related issues with programs that utilizes `malloc` . Specifically, it checks to see if you `free` all the memory you allocated.
2. To use this tool, type the following into the terminal

```

make program_name
valgrind ./program_name

```

Garbage Values

1. When asking the compiler for a block of memory, it is very likely that the memory has previously been utilized, so we might see *junk* or *garbage values*, as a result of asking for memory without initializing it.

Swapping

1. If we are calling `swap` function as followed, it fails to actually swap the value, because when passing values to a function, you are only providing copies. The *scope* of `x` and `y` is limited to the `main` function. `x` and `y` are being passed by *value*.

```
// Fails to swap two integers

#include <stdio.h>

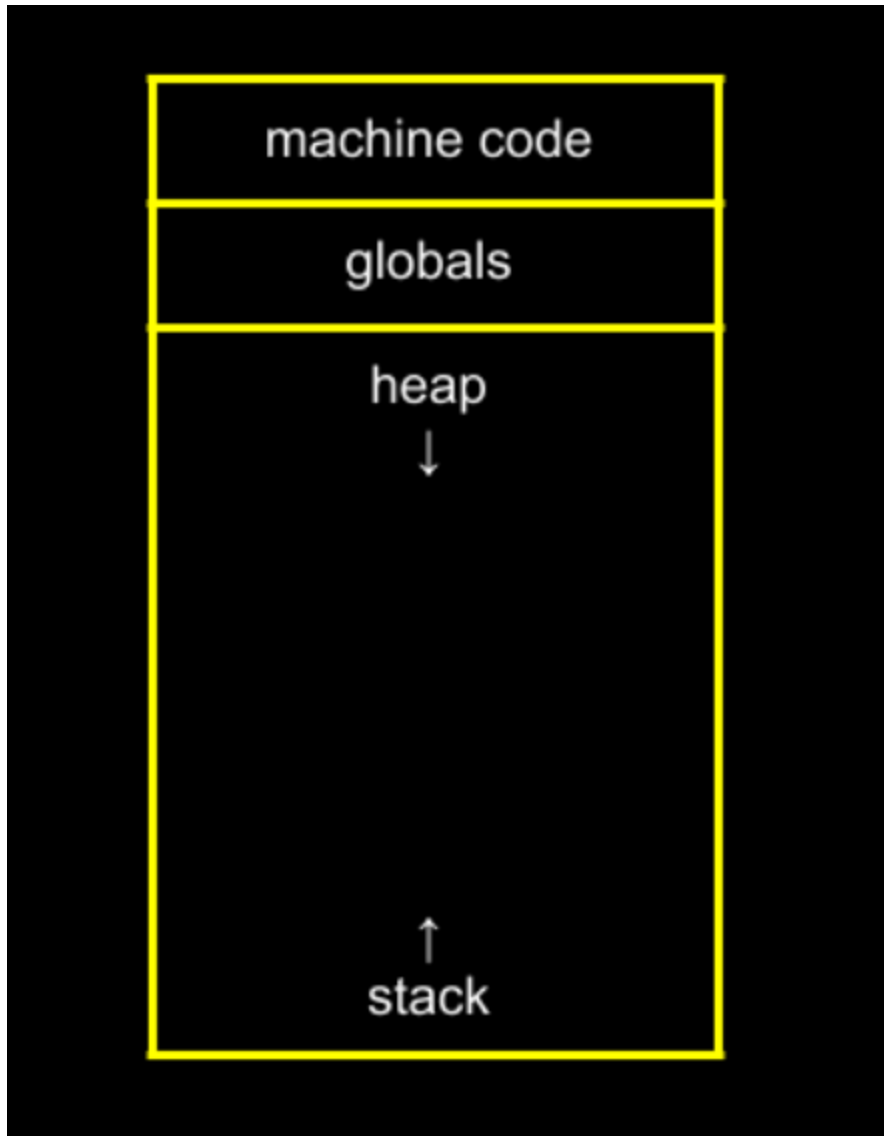
void swap(int a, int b);

int main(void)
{
    int x = 1;
    int y = 2;

    printf("x is %i, y is %i\n", x, y);
    swap(x, y);
    printf("x is %i, y is %i\n", x, y);
}

void swap(int a, int b)
{
    int tmp = a;
    a = b;
    b = tmp;
}
```


2. We can use the Memory Image to make this clearer:




Global variables live in one place in memory.

Various functions are stored in the `stack` in another area of memory.

3. The `main` function and the `swap` function is stored in the stack. They have two separate *frames* or areas of memory. Therefore, we cannot simply pass the values from one

function to another to change them.



4. So we use pointer to make the swap:

```
// Swaps two integers using pointers

#include <stdio.h>

void swap(int *a, int *b); // pass two pointers to the swap function, which are
the addresses of a and b

int main(void)
{
    int x = 1;
    int y = 2;

    printf("x is %i, y is %i\n", x, y);
    swap(&x, &y); // & is "the address of the variable"
    printf("x is %i, y is %i\n", x, y);
}
```


```
void swap(int *a, int *b)
{
    int tmp = *a; //create a tmp value that now has the value of
    x(dereferencing a means the value of the address that a stores, in other
    words, * means go to a and get the value, then pass it to tmp)
    *a = *b; // go to the address of a, find the value, replace the value of
    y(*b)
    *b = tmp; //go to the address of b, find the value, replace the value of
    x(tmp)
    // now the address of x contains the the value of y, vice versa, swap ends
}
```

Now variables are not passed by value but by reference. That is, the addresses of a and b are provided to the swap function.

Overflow

1. A *heap overflow* is when you overflow the heap, touching areas of memory you are not supposed to.
2. A *stack overflow* is when too many functions are called, overflowing the amount of memory available.

- Both of these are considered *buffer overflows*.



Scanf

- Replace `get_int` with `scanf`

```
printf("n: ");  
scanf("%i", &n);
```

where `%i` specifies the type of variable as input, `&n` specifies the address that the value is stored in

- Replace `get_string` with `scanf`

The problem is that we need to pre-allocate memory for the string to store it.

File I/O

1. Like in mps, sometimes we need to access information in another file like a `txt` file, so we need the ability to open a file in a `c` program.
2. CSV: comma-separated values
3. `fopen`

```
FILE *fopen(const char *filename, const char *mode); //prototype for fopen
```

Some modes:

"a": append mode: if the file exists, add the content at the end, else, create the file

"r": read mode: read the file with the file pointer initially pointing to the beginning of the file, if it DNE, fails and return `NULL`

An example:

```
// Saves names and numbers to a CSV file

#include <cs50.h>
#include <stdio.h>
#include <string.h>

int main(void)
{
    // Open CSV file
    FILE *file = fopen("phonebook.csv", "a");
    // FILE is a special data type, which is a file
    // The pointer named "file" points to a file
    // The fopen function opens a file named "phonebook.csv" in append mode
    if (!file)
    //if file is NULL, !file will be 1, if(1) will execute the return 1 inside
    {
        return 1;
    }

    // Get name and number
    char *name = get_string("Name: ");
    char *number = get_string("Number: ");

    // Print to file
    fprintf(file, "%s,%s\n", name, number);
    // name and number will be written into where file points to as strings

    // Close file
```

```
    fclose(file);  
}
```

Another example for copying a file to another file

```
// Copies a file  
  
#include <stdio.h>  
#include <stdint.h>  
  
typedef uint8_t BYTE;  
  
int main(int argc, char *argv[])  
// take in command-line argument, which contains names of the source file and  
// destination file  
{  
    FILE *src = fopen(argv[1], "rb");  
    // open the source file specified by the first command-line argument  
    // in binary read mode  
    FILE *dst = fopen(argv[2], "wb");  
    // open the destination file specified by the second command-line argument  
    // in binary write mode  
  
    BYTE b;  
    // the copy process hope to copy byte by byte, which is why we use rb and  
wb  
  
    while (fread(&b, sizeof(b), 1, src) != 0)  
    // read from file src, read 1 sizeof(b) and stores in b every time  
    {  
        fwrite(&b, sizeof(b), 1, dst);  
        // write 1 sizeof(b) from b into dst file every time  
    }  
  
    fclose(dst);  
    fclose(src);  
}
```

Takeaways

1. One thing about array is that the name of the array itself is the address of the first entry of the array, i.e.

```
char s[4]; //s=&s[0]
```